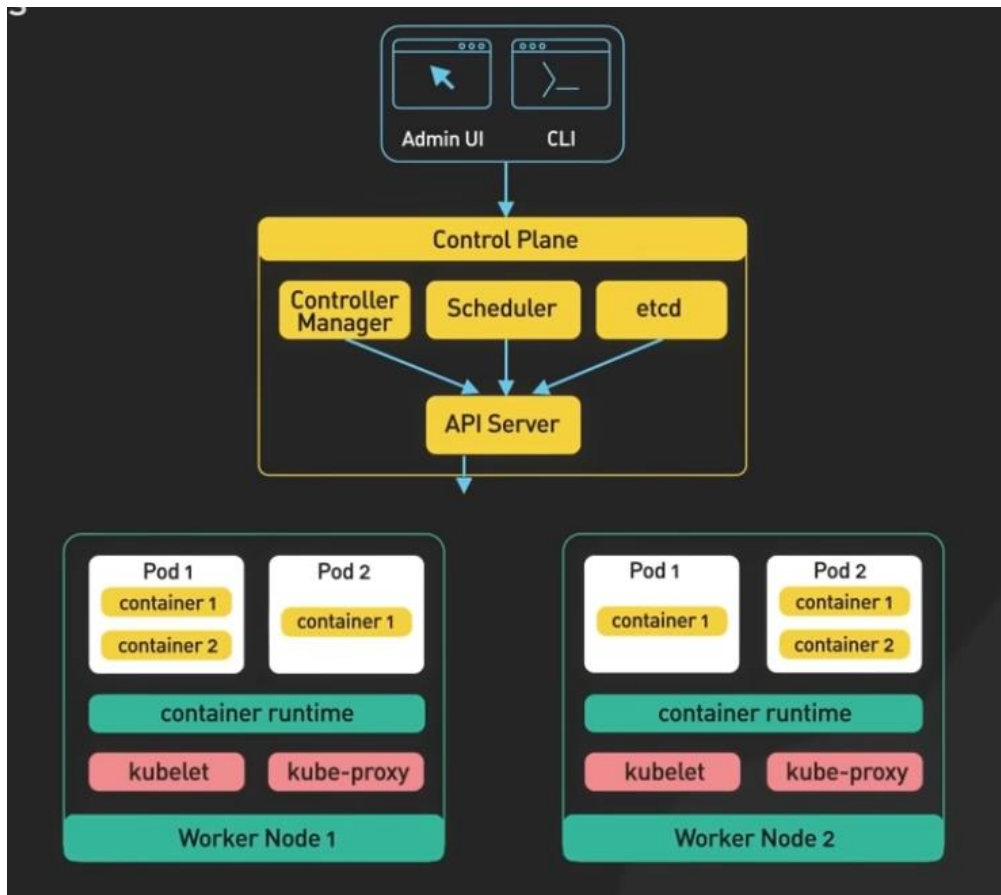


## Control Plane:



### 1. API Server (**kube-apiserver**)

- **Role:** Frontend REST API for cluster operations.
- **Key Commands:**
  - Check status:  
`kubectl get componentstatuses (kubectl get cs)`
  - View logs (systemd):  
`journalctl -u kube-apiserver`
  - Health endpoint:  
`curl -k https://<API_SERVER_IP>:6443/healthz`

### 2. etcd

- **Role:** Distributed key-value store for cluster state.
- **Key Commands:**
  - Check cluster health:  
`etcdctl endpoint health --endpoints=<ETCD_ENDPOINTS>`
  - List members:  
`etcdctl member list`
  - Snapshot backup:  
`etcdctl snapshot save /path/to/snapshot.db`

- Restore snapshot:  
`etcdctl snapshot restore /path/to/snapshot.db`

### 3. Scheduler (**kube-scheduler**)

- **Role:** Assigns pods to nodes.
- **Key Commands:**
  - View logs:  
`journalctl -u kube-scheduler`
  - Check scheduling decisions:  
`kubectl describe pod <POD_NAME> | grep Events`

### 4. Controller Manager (**kube-controller-manager**)

- **Role:** Runs core controllers (node, replication, endpoints).
- **Key Commands:**
  - View logs:  
`journalctl -u kube-controller-manager`

### 6. Kubelet

- **Role:** Runs on nodes, manages pods and containers.
- **Key Commands:**
  - Check status:  
`systemctl status kubelet`
  - View logs:  
`journalctl -u kubelet`

### 7. Kubectl (CLI Tool)

- **Common Commands:**
  - **Get resources:** `kubectl get <pods/nodes/services>`
  - **Describe resource:** `kubectl describe <pod/node> <NAME>`
  - **View logs:** `kubectl logs <POD_NAME>`
  - **Exec into pod:** `kubectl exec -it <POD_NAME> -- /bin/sh`

### 8. kube-proxy

it is responsible for managing network rules to ensure communication between **services**, **pods**, and external clients works correctly.

#### Key Commands

- **Check Service Endpoints:**  
`kubectl get endpoints <SERVICE_NAME>`
- **Check kube-proxy config:**  
`kubectl describe configmap -n kube-system kube-proxy`

## 11. Namespaces

- **default**: Created automatically.
- **kube-system**: Hosts Control Plane components (e.g., **kube-proxy**, **CoreDNS**).
- **kube-public**: Readable by all users (rarely used).
- **kube-node-lease**: Node heartbeat tracking.

```
Namespace.yaml
apiVersion: v1
kind: Namespace
metadata:
  name: dev
---
apiVersion: v1
kind: Namespace
metadata:
  name: prod
```

```
12. Pod
apiVersion: v1
kind: Pod
metadata:
  name: advanced-pod
  labels:
    app: nginx
  namespace: production
  tier: frontend
spec:
  # --- Containers Section ---
  containers:
  - name: nginx
    image: nginx:1.25-alpine
    imagePullPolicy: IfNotPresent # Options: Always/Never/IfNotPresent

  # --- Ports ---
  ports:
  - containerPort: 80
    protocol: TCP
    name: http

  # --- Resource Limits ---
  resources:
    requests:
      cpu: "250m" # 0.25 CPU cores
      memory: "256Mi" # 256 MB RAM
    limits:
```

cpu: "1" # 1 CPU core max  
memory: "512Mi" # 512 MB RAM max

# --- Liveness Probe ---

livenessProbe:

httpGet:

path: /healthz

port: 80

initialDelaySeconds: 15 # Wait 15s before first probe

periodSeconds: 10 # Check every 10s

timeoutSeconds: 2 # Probe timeout

failureThreshold: 3 # 3 failures = container restart

# --- Readiness Probe ---

readinessProbe:

tcpSocket:

port: 80

initialDelaySeconds: 5

periodSeconds: 5

successThreshold: 1

failureThreshold: 3

# --- Environment Variables ---

env:

- name: NGINX\_ENV

value: "production"

- name: DB\_HOST

valueFrom:

secretKeyRef:

name: db-secret

key: host

# --- Volume Mounts ---

volumeMounts:

- name: config-volume

mountPath: /etc/nginx/conf.d

readOnly: true

- name: logs-volume

mountPath: /var/log/nginx

# --- Security Context ---

securityContext:

runAsNonRoot: true

runAsUser: 1000

runAsGroup: 3000

```
fsGroup: 2000
readOnlyRootFilesystem: true
capabilities:
  drop: ["ALL"]
  add: ["NET_BIND_SERVICE"]
```

```
# --- Init Containers ---
```

```
initContainers:
- name: init-db
  image: busybox:1.28
  command: ['sh', '-c', 'until nslookup db-service; do echo waiting for db; sleep 2; done']
```

```
# --- Volumes ---
```

```
volumes:
- name: config-volume
  configMap:
    name: nginx-config
- name: logs-volume
  emptyDir: {}
```

```
# --- Pod-wide Settings ---
```

```
restartPolicy: Always # Options: Always/OnFailure/Never
nodeSelector:
  disktype: ssd # Schedule on nodes with this label
tolerations:
- key: "special"
  operator: "Exists"
  effect: "NoSchedule"
securityContext:
  fsGroup: 2000 # File ownership group
```

```
Deployments
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
    tier: frontend
spec:
  replicas: 3 # Number of Pods to maintain
  selector:
    matchLabels:
      app: nginx
      tier: frontend
```

```
strategy:
  type: RollingUpdate
  rollingUpdate:
    maxSurge: 1      # How many extra Pods can be created during update
    maxUnavailable: 0 # No Pods should be unavailable during update
template:
  metadata:
    labels:
      app: nginx
      tier: frontend
  spec:
    # --- Containers Section ---
    containers:
      - name: nginx
        image: nginx:1.25-alpine
        imagePullPolicy: IfNotPresent
        ports:
          - containerPort: 80
            protocol: TCP
            name: http
        resources:
          requests:
            cpu: "250m"
            memory: "256Mi"
          limits:
            cpu: "1"
            memory: "512Mi"
        livenessProbe:
          httpGet:
            path: /healthz
            port: 80
          initialDelaySeconds: 15
          periodSeconds: 10
          timeoutSeconds: 2
          failureThreshold: 3
        readinessProbe:
          tcpSocket:
            port: 80
          initialDelaySeconds: 5
          periodSeconds: 5
          successThreshold: 1
          failureThreshold: 3
        env:
          - name: NGINX_ENV
            value: "production"
```

```
- name: DB_HOST
valueFrom:
  secretKeyRef:
    name: db-secret
    key: host
volumeMounts:
- name: config-volume
  mountPath: /etc/nginx/conf.d
  readOnly: true
- name: logs-volume
  mountPath: /var/log/nginx
securityContext:
  runAsNonRoot: true
  runAsUser: 1000
  readOnlyRootFilesystem: true
  capabilities:
    drop: ["ALL"]
    add: ["NET_BIND_SERVICE"]
```

# --- Init Containers ---

```
initContainers:
- name: init-db
  image: busybox:1.28
  command: ['sh', '-c', 'until nslookup db-service; do echo waiting for db; sleep 2; done']
```

# --- Volumes ---

```
volumes:
- name: config-volume
  configMap:
    name: nginx-config
- name: logs-volume
  emptyDir: {}
```

# --- Pod-wide Settings ---

```
nodeSelector:
  disktype: ssd
tolerations:
- key: "special"
  operator: "Exists"
  effect: "NoSchedule"
securityContext:
  fsGroup: 2000
```

```
Replica set
apiVersion: apps/v1
kind: ReplicaSet
```

```
metadata:
  name: frontend-rs
  labels:
    app: frontend
    tier: web
spec:
  replicas: 3
  selector:
    matchLabels:
      tier: web # Must match pod template labels
  template:
    metadata:
      labels:
        tier: web # Must match selector
    spec:
      containers:
        - name: nginx
          image: nginx:1.25
          ports:
            - containerPort: 80
```

## SERVICES

### 1. ClusterIP (Internal Service)

```
apiVersion: v1
kind: Service
metadata:
  name: nginx-service
  labels:
    app: nginx
    tier: frontend
spec:
  type: ClusterIP # Default - accessible only within the cluster
  selector:
    app: nginx # Must match Deployment's pod labels
    tier: frontend
  ports:
    - name: http
      port: 80 # Service port
      targetPort: 80 # Pod port (matches containerPort)
      protocol: TCP
```

### 2. NodePort (External Access via Node IP)

```
apiVersion: v1
kind: Service
metadata:
```

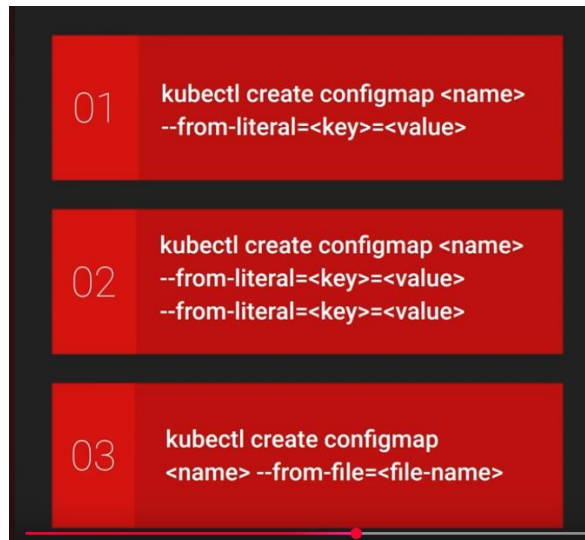
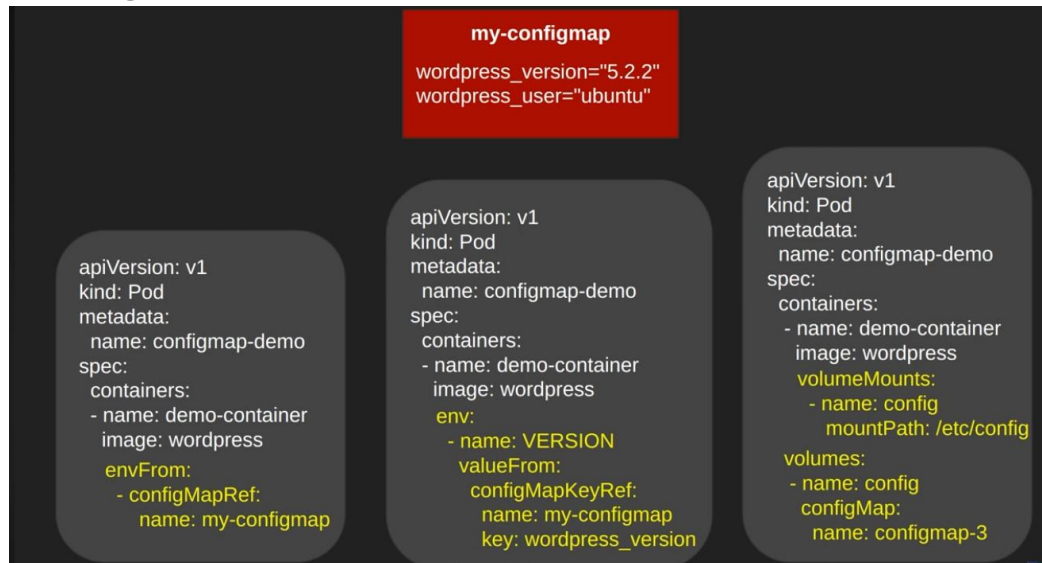


```
name: nginx-nodeport
spec:
  type: NodePort
  selector:
    app: nginx
    tier: frontend
  ports:
  - port: 80
    targetPort: 80
    nodePort: 30080 # Optional (default range: 30000-32767)
```

### 3. LoadBalancer (Cloud Provider Integration)

```
apiVersion: v1
kind: Service
metadata:
  name: nginx-lb
  annotations:
    cloud-provider-specific-key: "value" # e.g., AWS: service.beta.kubernetes.io/aws-load-
balancer-type: nlb
spec:
  type: LoadBalancer
  selector:
    app: nginx
    tier: frontend
  ports:
  - port: 80
    targetPort: 80
  externalTrafficPolicy: Local # Preserves client IP
```

# Config Map



## Secrets

### Generic Secret

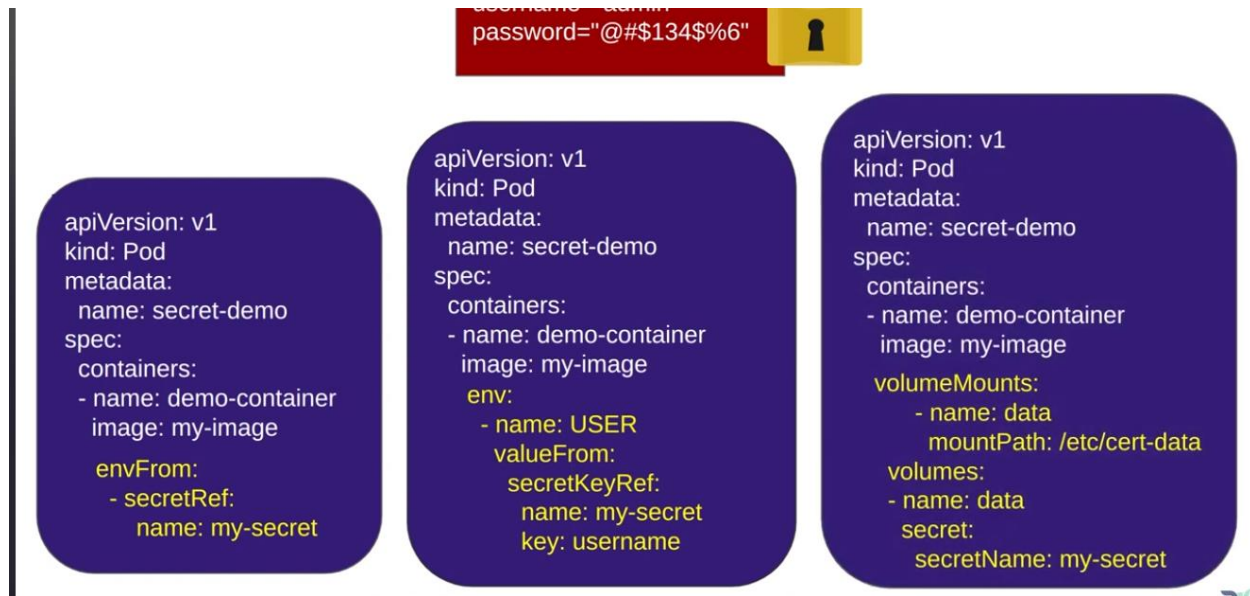
`kubectl create secret generic db-secret --from-literal=username=dbuser --from-literal=password=Y4nys7f11`

### Docker-registry Secret

`kubectl create secret docker-registry docker-secret --docker-email=example@gmail.com --docker-username=dev --docker-password=pass1234 --docker-server=my-registry.example:5000`

### TLS Secret

`kubectl create secret tls my-tls-secret --cert=/root/data/serverca.crt --key=/root/data/servercakey.pem`



## Volumes, PV,PVC

### Workflow

1. **Admin** creates StorageClass (defines storage types).
2. **User** creates pvc (requests storage).
3. **Kubernetes:**
  - o Dynamically provisions PV (cloud) **or**
  - o Binds to existing PV (on-prem)
4. **Pod** mounts PVC → accesses persistent storage.

### Storage class

```

apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
  name: ebs-sc
provisioner: ebs.csi.aws.com
parameters:
  type: gp2
reclaimPolicy: Delete # Or 'Retain'
volumeBindingMode: WaitForFirstConsumer
  
```

### pvc.yaml

```

apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: ebs-claim
spec:
  accessModes:
    - ReadWriteOnce # RWO, ROX, RWX
  
```

```
storageClassName: ebs-sc
resources:
  requests:
    storage: 1Gi
```

### **pod.yaml**

```
//Baki file of pod
  volumeMounts:
    - name: persistent-storage
      mountPath: /data
  volumes:
    - name: persistent-storage
      persistentVolumeClaim:
        claimName: ebs-claim
```

### Service Account

#### Role.yaml

```
apiVersion: rbac.authorization.k8s.io/v1
kind: Role
metadata:
  namespace: default
  name: pod-reader
rules:
- apiGroups:
  - ""
  resources:
  - pods
  verbs:
  - get
  - watch
  - list
```

#### RoleBinding.yaml

```
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: read-pods
  namespace: default
subjects:
- kind: ServiceAccount
  name: mysa
  namespace: default
roleRef:
  kind: Role
  name: pod-reader
  apiGroup: rbac.authorization.k8s.io
```

In pod.yaml

spec:

serviceAccountName: mysa

Verifying Permissions => kubectl auth can-i get pods --as=system:serviceaccount:default:mysa

## HELM

my-chart/

- ├── Chart.yaml # Chart metadata
- ├── values.yaml # Default configuration
- ├── charts/ # Sub-charts
- ├── templates/ # Kubernetes manifests
  - ├── deployment.yaml
  - ├── service.yaml
  - └── \_helpers.tpl # Reusable templates
- └── templates/NOTES.txt # Post-install notes

templates/deployment.yaml

apiVersion: apps/v1

kind: Deployment

metadata:

name: {{ .Release.Name }}-app # Built-in object

spec:

replicas: {{ .Values.replicaCount }} # Value reference

template:

spec:

containers:

- name: {{ .Chart.Name }} # Built-in

image: "{{ .Values.image.repo }}:{{ .Values.image.tag }}"

# --- IF/ELSE EXAMPLE ---

{{- if .Values.resources }}

resources: {{ toYaml .Values.resources | nindent 10 }}

{{- else }}

resources: {}

{{- end }}

# --- RANGE EXAMPLE ---

env:

{{- range \$key, \$val := .Values.env }}

- name: {{ \$key | upper }}

value: {{ \$val | quote }}

{{- end }}

# --- WITH SCOPE EXAMPLE ---

{{- with .Values.securityContext }}

securityContext: {{ . | toYaml | nindent 10 }}

{{- end }}

values.yaml

# --- BASIC VALUES ---

replicaCount: 2

# --- NESTED VALUES ---

image:

repo: nginx

tag: latest

# --- IF/ELSE TARGET ---

resources:

limits:

cpu: "500m"

memory: "512Mi"

# --- RANGE TARGET ---

env:

log\_level: debug

cache\_enabled: "true"

# --- WITH SCOPE TARGET ---

securityContext:

runAsNonRoot: true

{{ .Values.text | quote }} # Add quotes

{{ .Values.data | toYaml | indent 4 }} # YAML+indent

# Install with overrides

helm install my-app . --set replicaCount=3

# Dry-run template rendering

helm template . --set image.tag=canary

Named Template

\_helpers.tpl

{{- define "labels" }}

app: nginx

env: dev

{{- end }}

Deployment.yaml

apiVersion: apps/v1

kind: Deployment

metadata:

```
name: nginx-deployment
labels:
  {{- template "labels" }}
spec:
  replicas: 1
  selector:
    matchLabels:
      {{- include "labels" . | indent 2 }}
  template:
    metadata:
      labels:
        {{- include "labels" . | indent 4 }}
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 80
```